

Programming for SAS Viya Users

Research Triangle SAS Users Group

20 June 2024

Charu Shankar
SAS Institute Canada




Bio

With a background in computer systems management. SAS Instructor Charu Shankar engages with logic, visuals, and analogies to spark critical thinking.

A SAS Instructor since 2007, she curates and delivers unique content via the SAS YouTube channel, SAS global forum, SAS Ask the Expert Series, SAS Training Post Blog, etc.

Charu loves to support users by teaching at conferences on topics related to SAS, SQL, Efficiencies, PERL, Macros, Python, Viya, etc.

When she's not coding, Charu is A Yoga Instructor who loves to explore Canadian trails with her husky Miko.



Agenda



What's SAS Viya - In Memory



Engines



Applications



Data – Types, Database Connections



Code

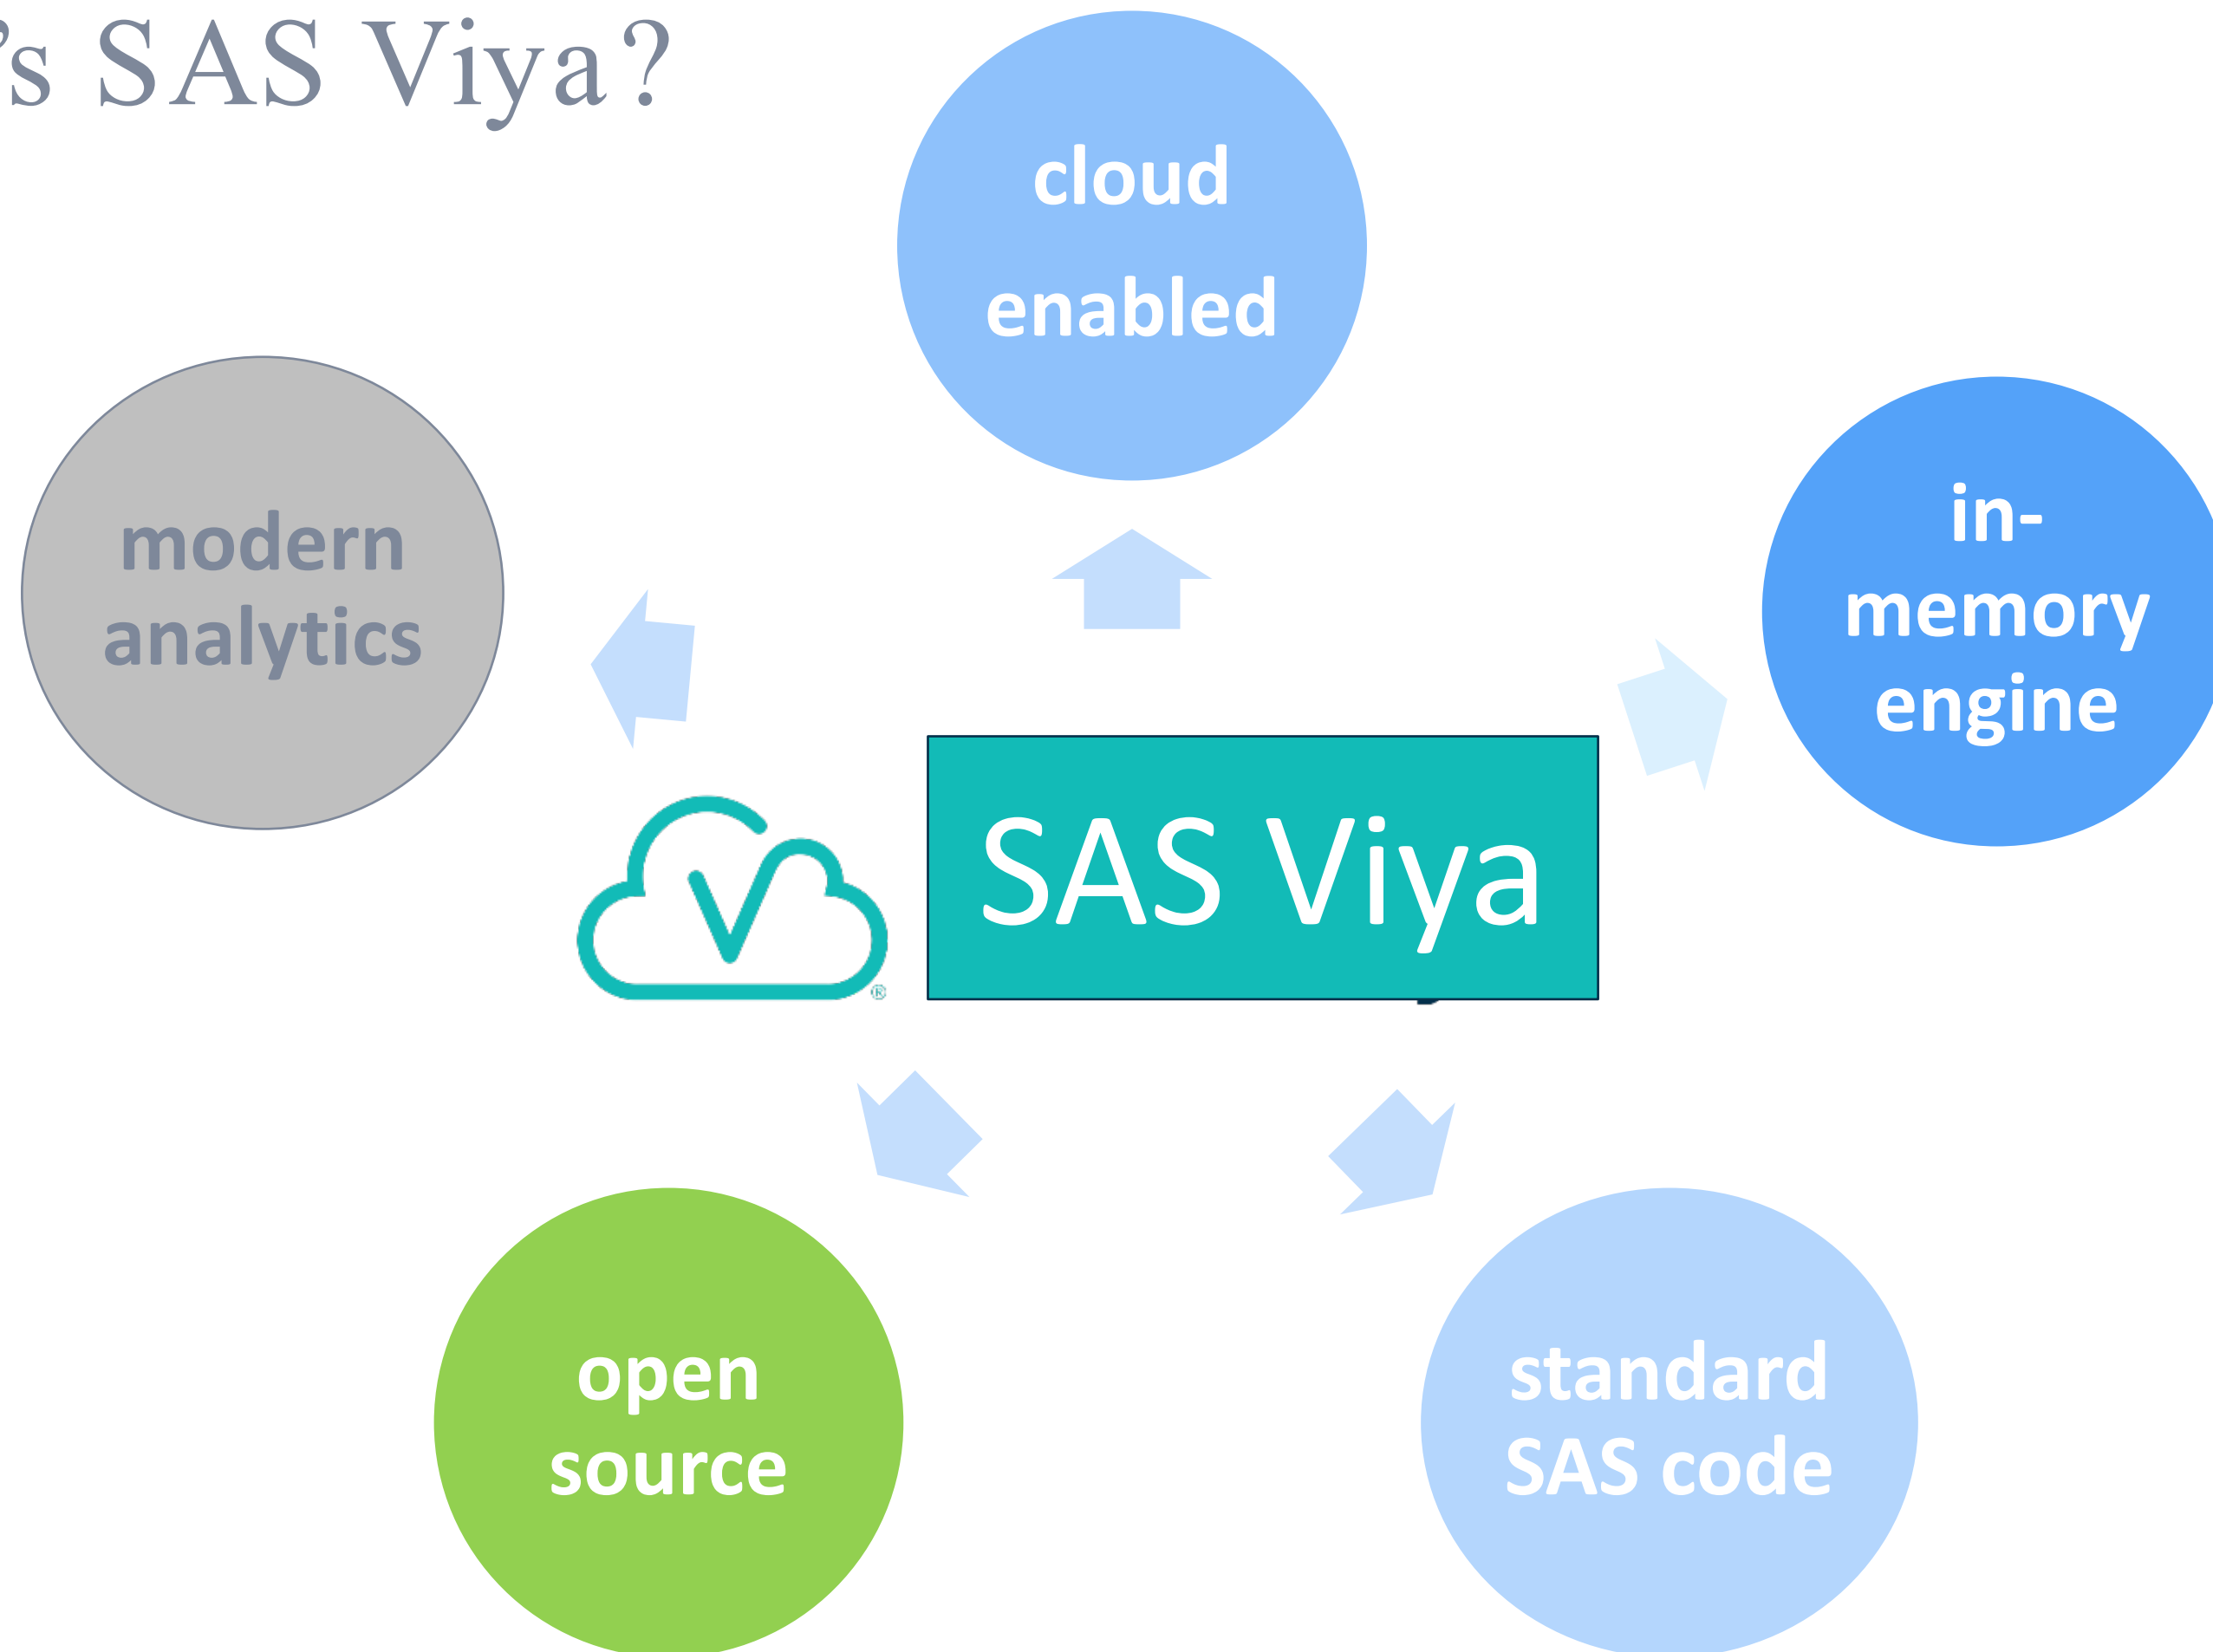


Demo



Handy Links

1 What's SAS Viya?



2 SAS Viya Engines


SAS Viya



Traditional SAS
processing engine

Executes traditional
SAS®9 code

SAS
Compute
Server



Next-gen SAS
processing engine

SAS Cloud
Analytic
Services
(CAS)

Executes CAS-enabled
code in parallel on
in-memory data



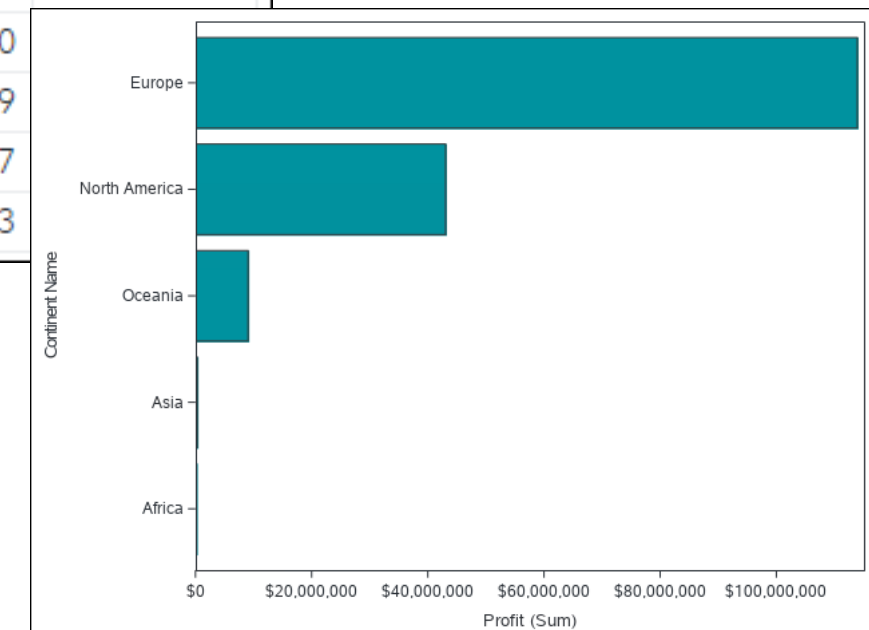
SAS Viya Servers

SAS
Compute
Server

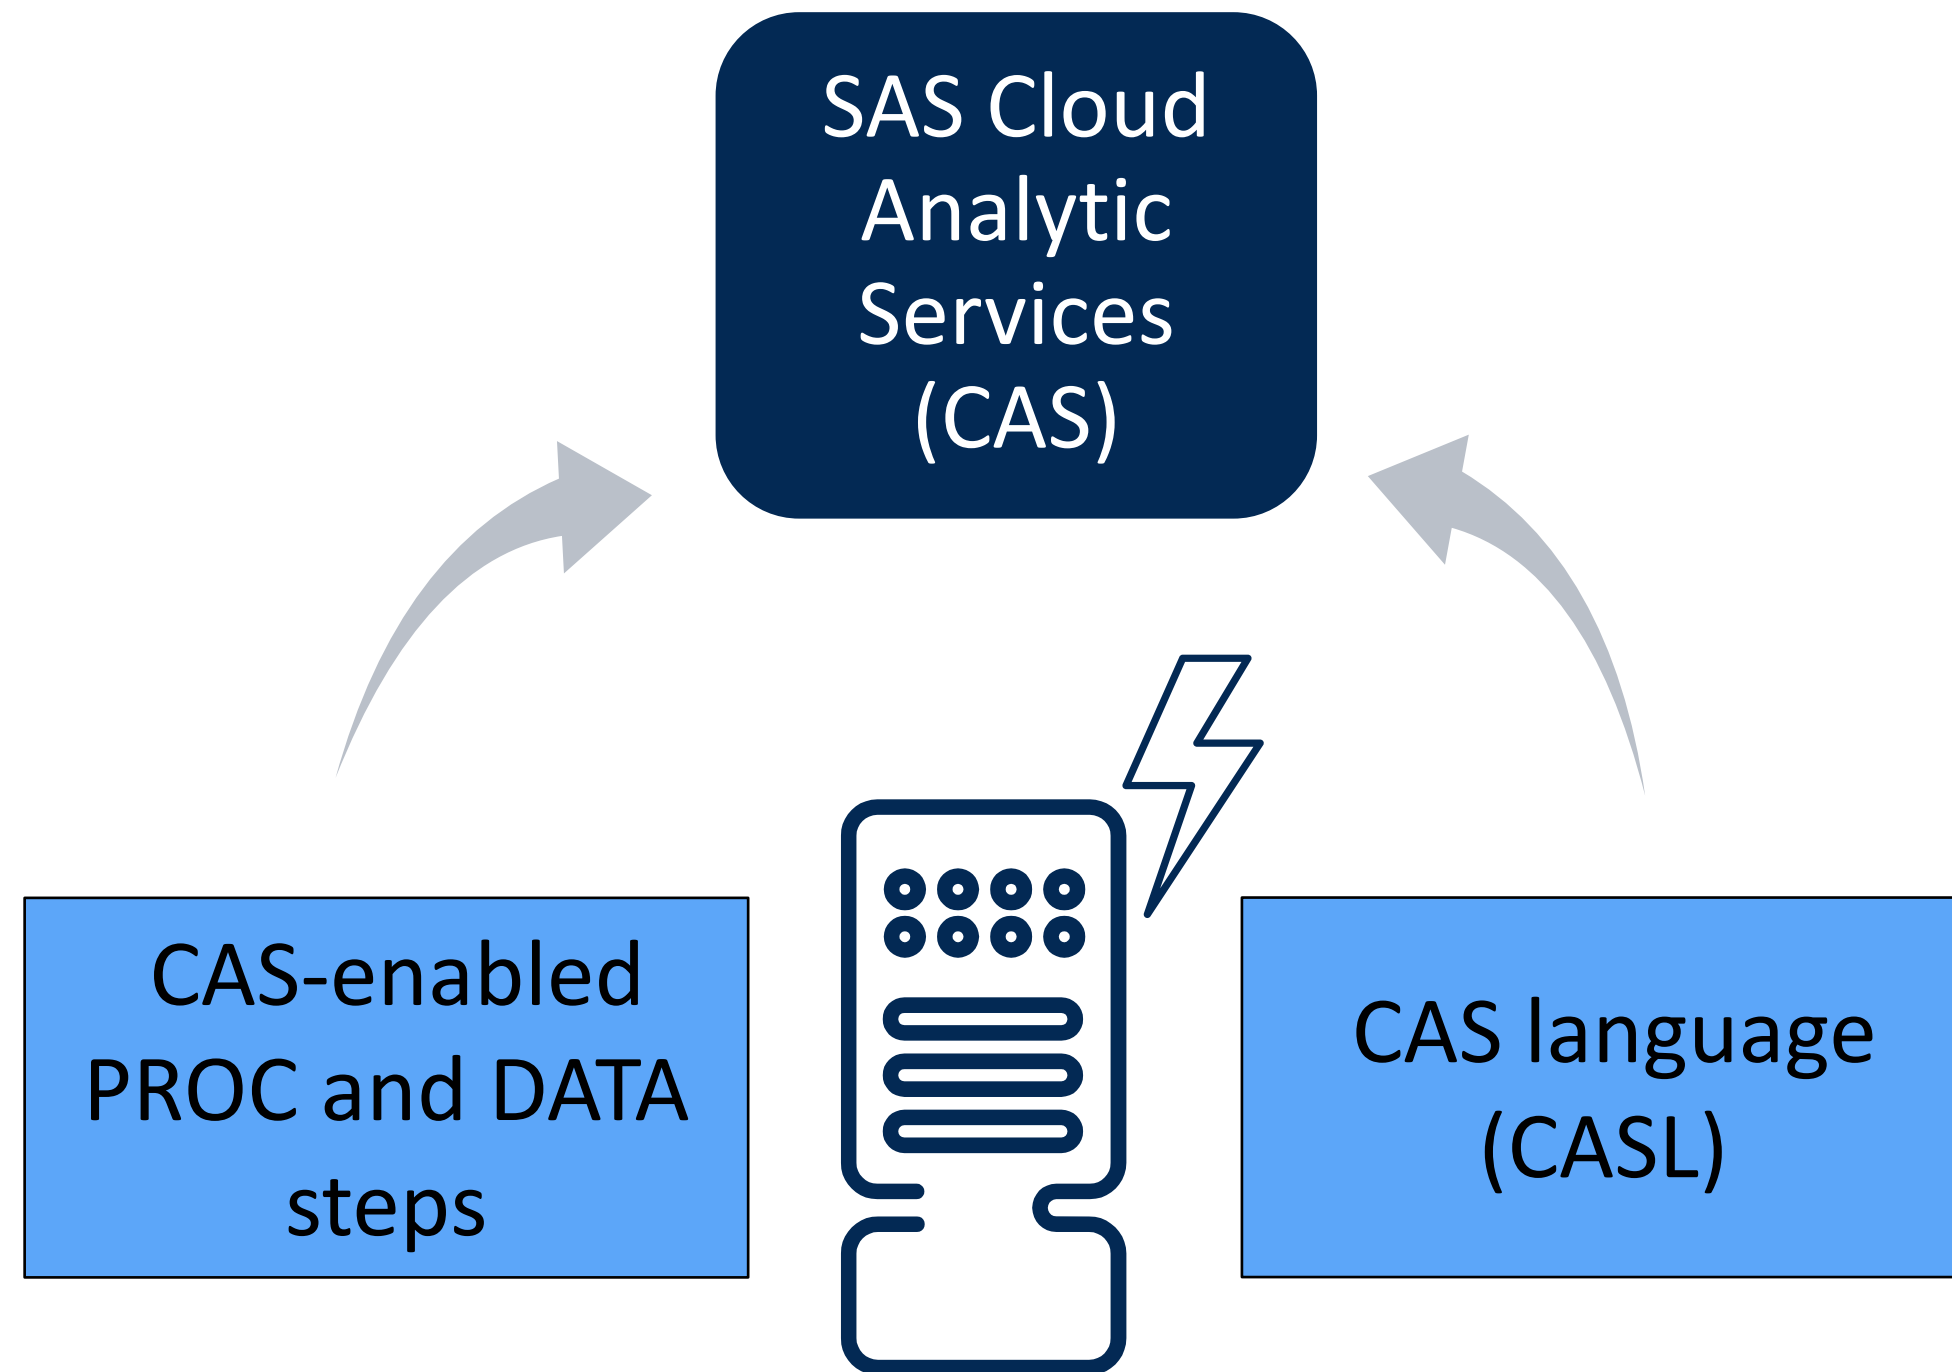
Standard SAS code
executes on the SAS
Compute Server.

```
libname pvbase "&path/data";  
  
data profit;  
    set pvbase.orders;  
    ...  
run;  
  
proc means data=profit;  
    ...  
run;
```


The MEANS Procedure			
Analysis Variable : Profit			
Continent Name	N Obs	Sum	Mean
Africa	770	-127.6800000	-0.1658182
Asia	1110	15503.70	
Europe	653684	5659450.59	
North America	235708	2121645.57	
Oceania	60397	462934.63	



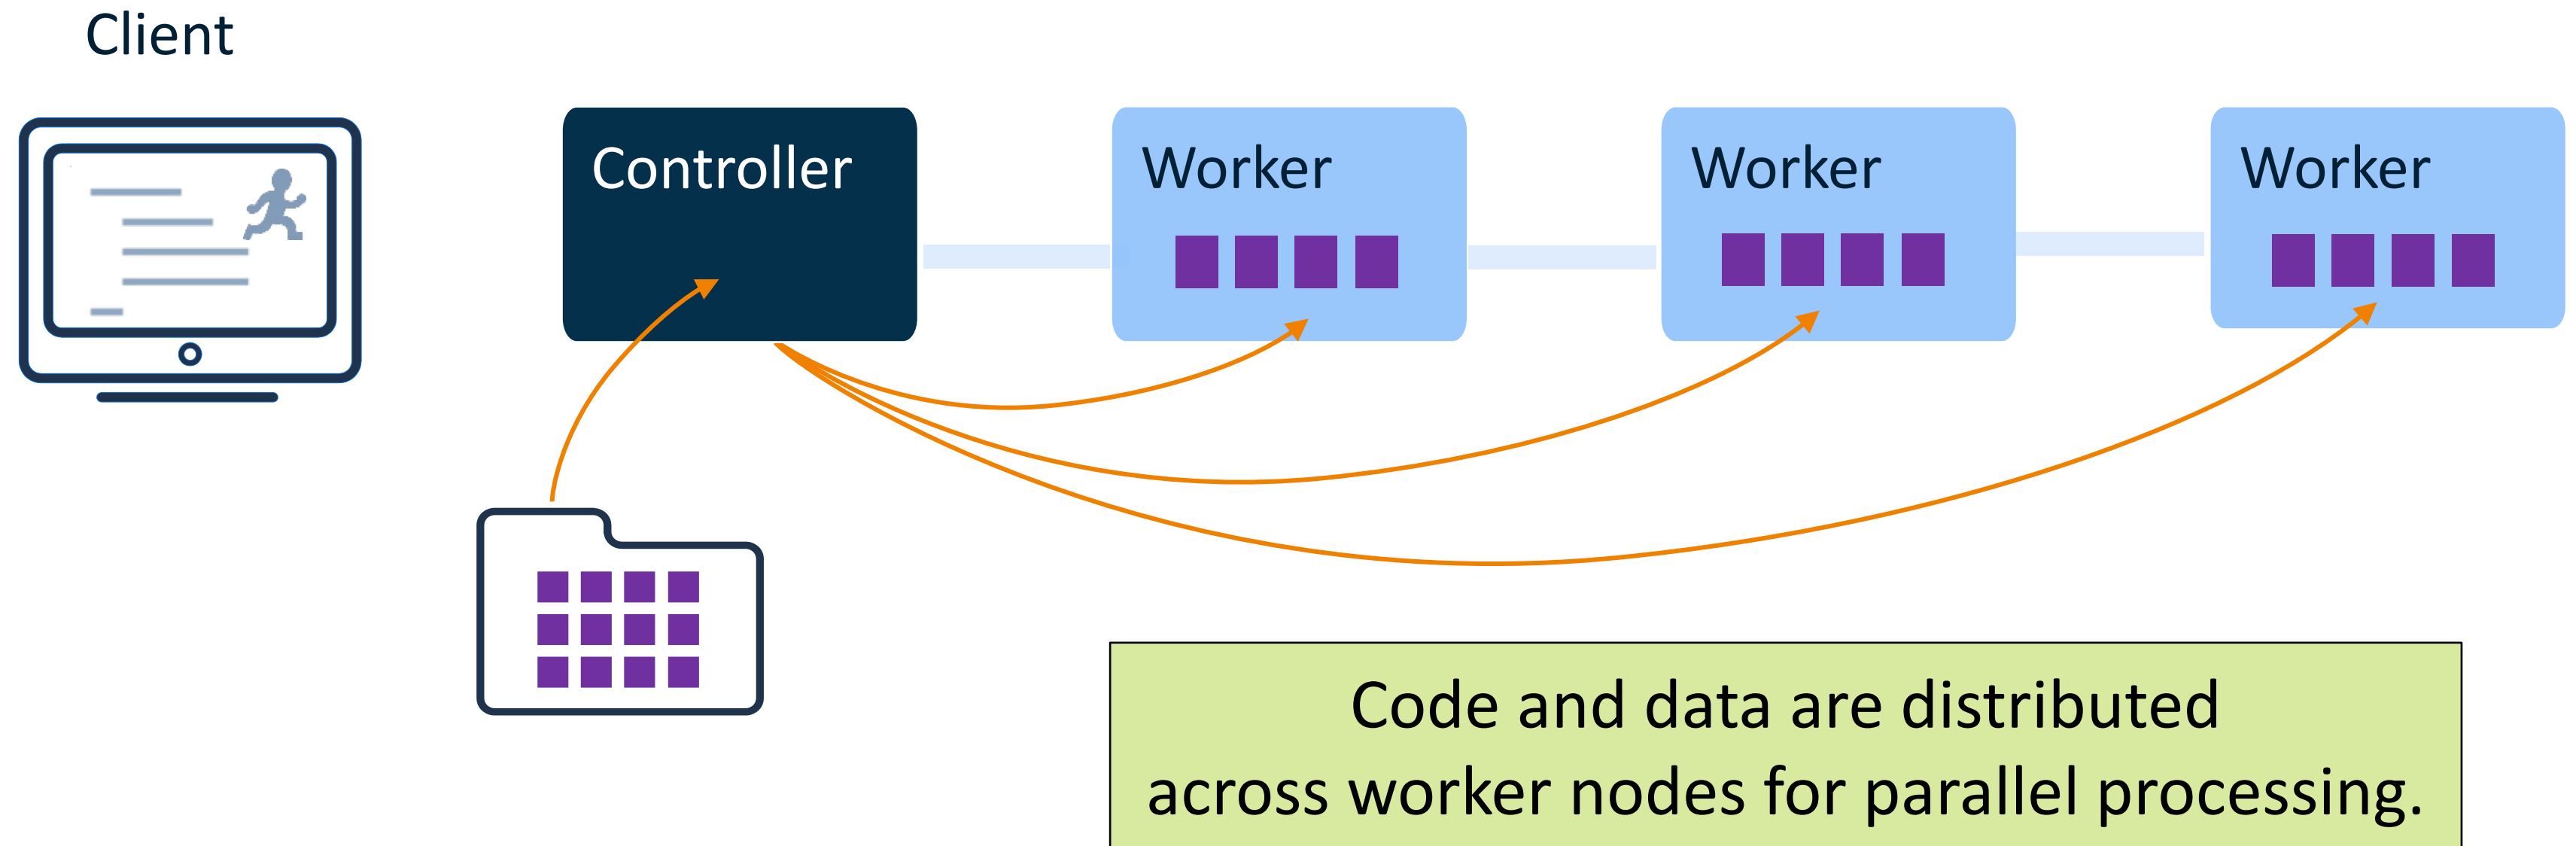
SAS Viya Servers



There are different ways to execute code in CAS.



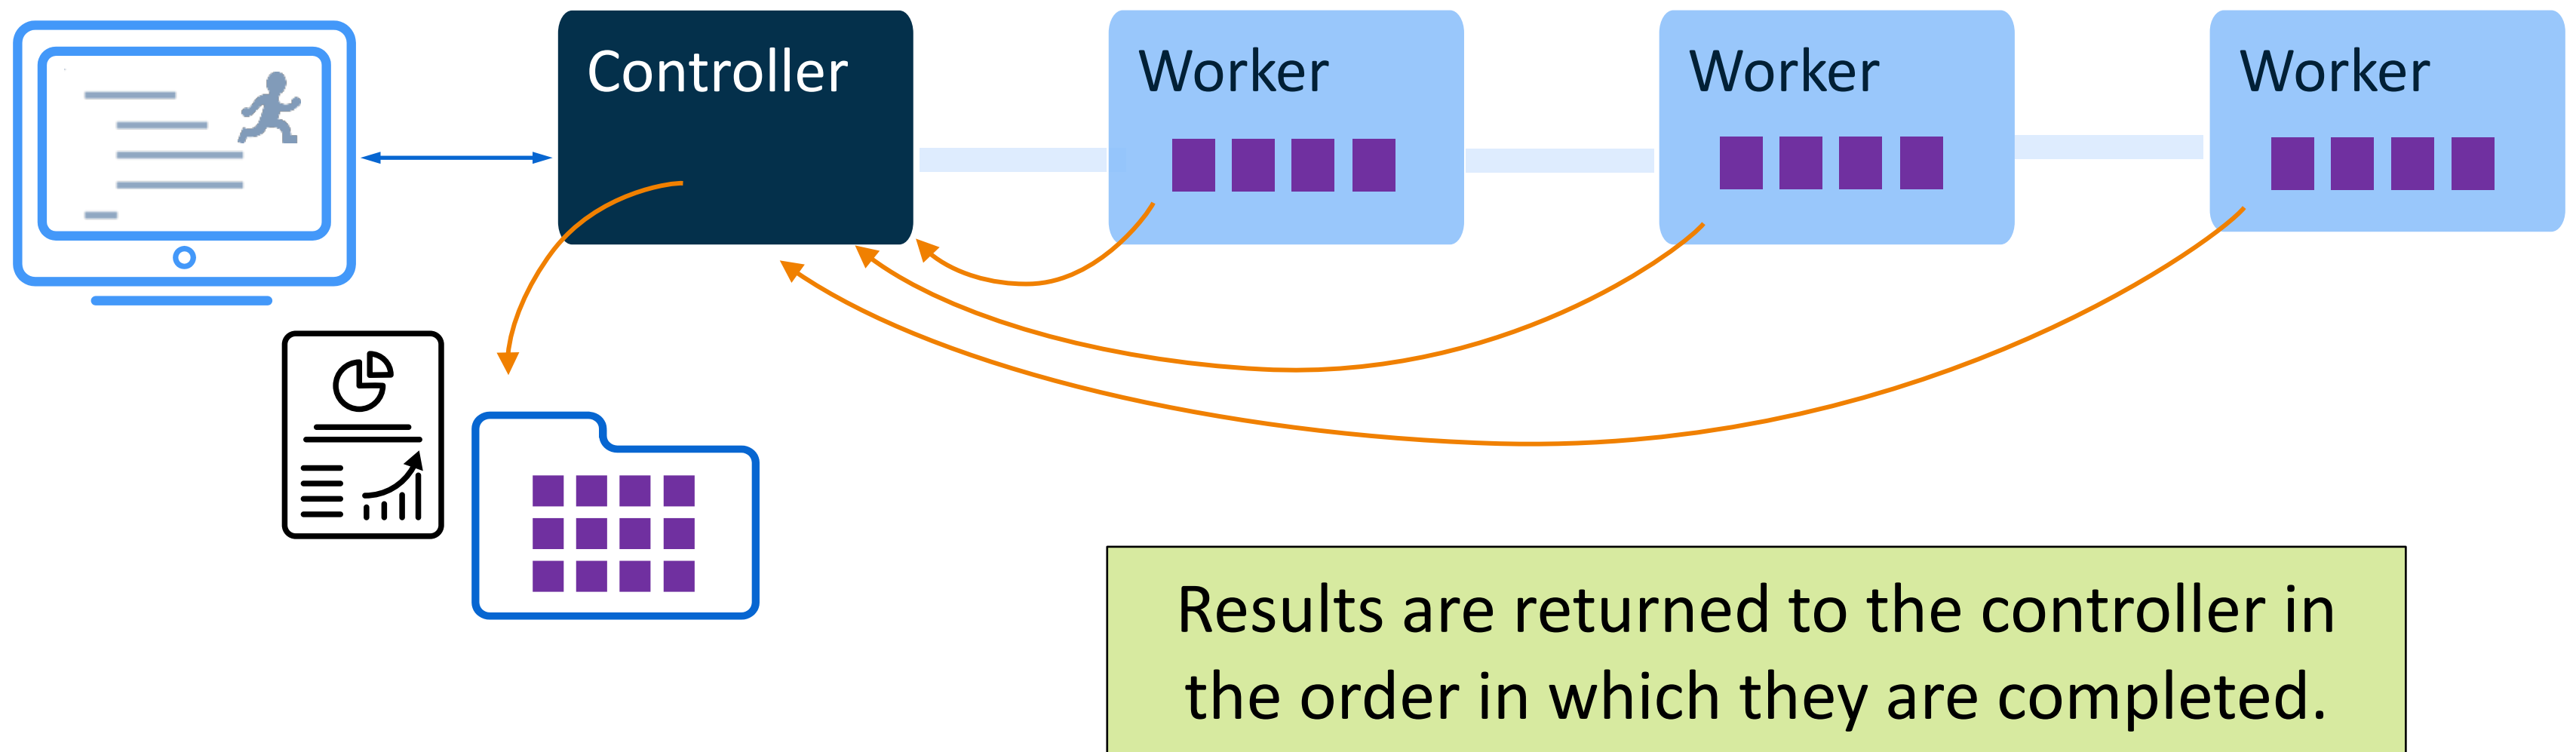
Parallel Processing in CAS



Parallel Processing in the Real World



Parallel Processing in CAS



Parallel Processing in the Real World



3 Applications

SAS Studio



Applications

SAS Visual Analytics



SAS Visual Data Mining
and Machine Learning



SAS Viya Programming Interface

The screenshot displays the SAS Studio web interface. The top navigation bar includes 'New', 'Options', 'View', 'Open', and 'Save All'. The left sidebar shows a 'Libraries' pane with a tree view containing 'MAPS', 'MAPSGFK', 'MAPSSAS', 'PV', and 'WORK'. The main editor area shows a SAS code file named 'pv01d01.sas' with the following code:

```
1 %let homedir=%sysget(HOME);
2 %let path=&homedir/Courses/PGVY35;
3
4 libname pv "&path/data";
5
6 data profit;
7   set pv.customers;
8   Profit=(RetailPrice-Cost)*Quantity;
9   format Profit dollar8.;
10 run;
11
12 ods excel file="&path/output/customers.xlsx";
13 proc means data=profit sum mean;
14   var Profit;
15   class Continent;
16 run;
17
18 proc sgplot data=profit;
19   hbar Continent / response=Profit stat=sum
20     categoryorder=respdesc;
21 run;
22 ods excel close;
```

The right pane shows the 'Results' tab with a table titled 'The MEANS Procedure' and a bar chart titled 'Analysis Variates'. The table displays the number of observations (N Obs) for each continent. The bar chart shows the total profit for each continent, with Europe having the highest profit.

Continent Name	N Obs
Africa	770
Asia	1110
Europe	653684
North America	235708
Oceania	60397

The bar chart shows the total profit for each continent, with Europe having the highest profit, followed by North America, and then Oceania.

SAS Studio

- web-based application
- access data and programs
- write new programs
- store code for common actions
- generate code with tasks

Other SAS Viya Programming Interfaces

Open-Source Applications

- Jupyter Notebook
- R Studio

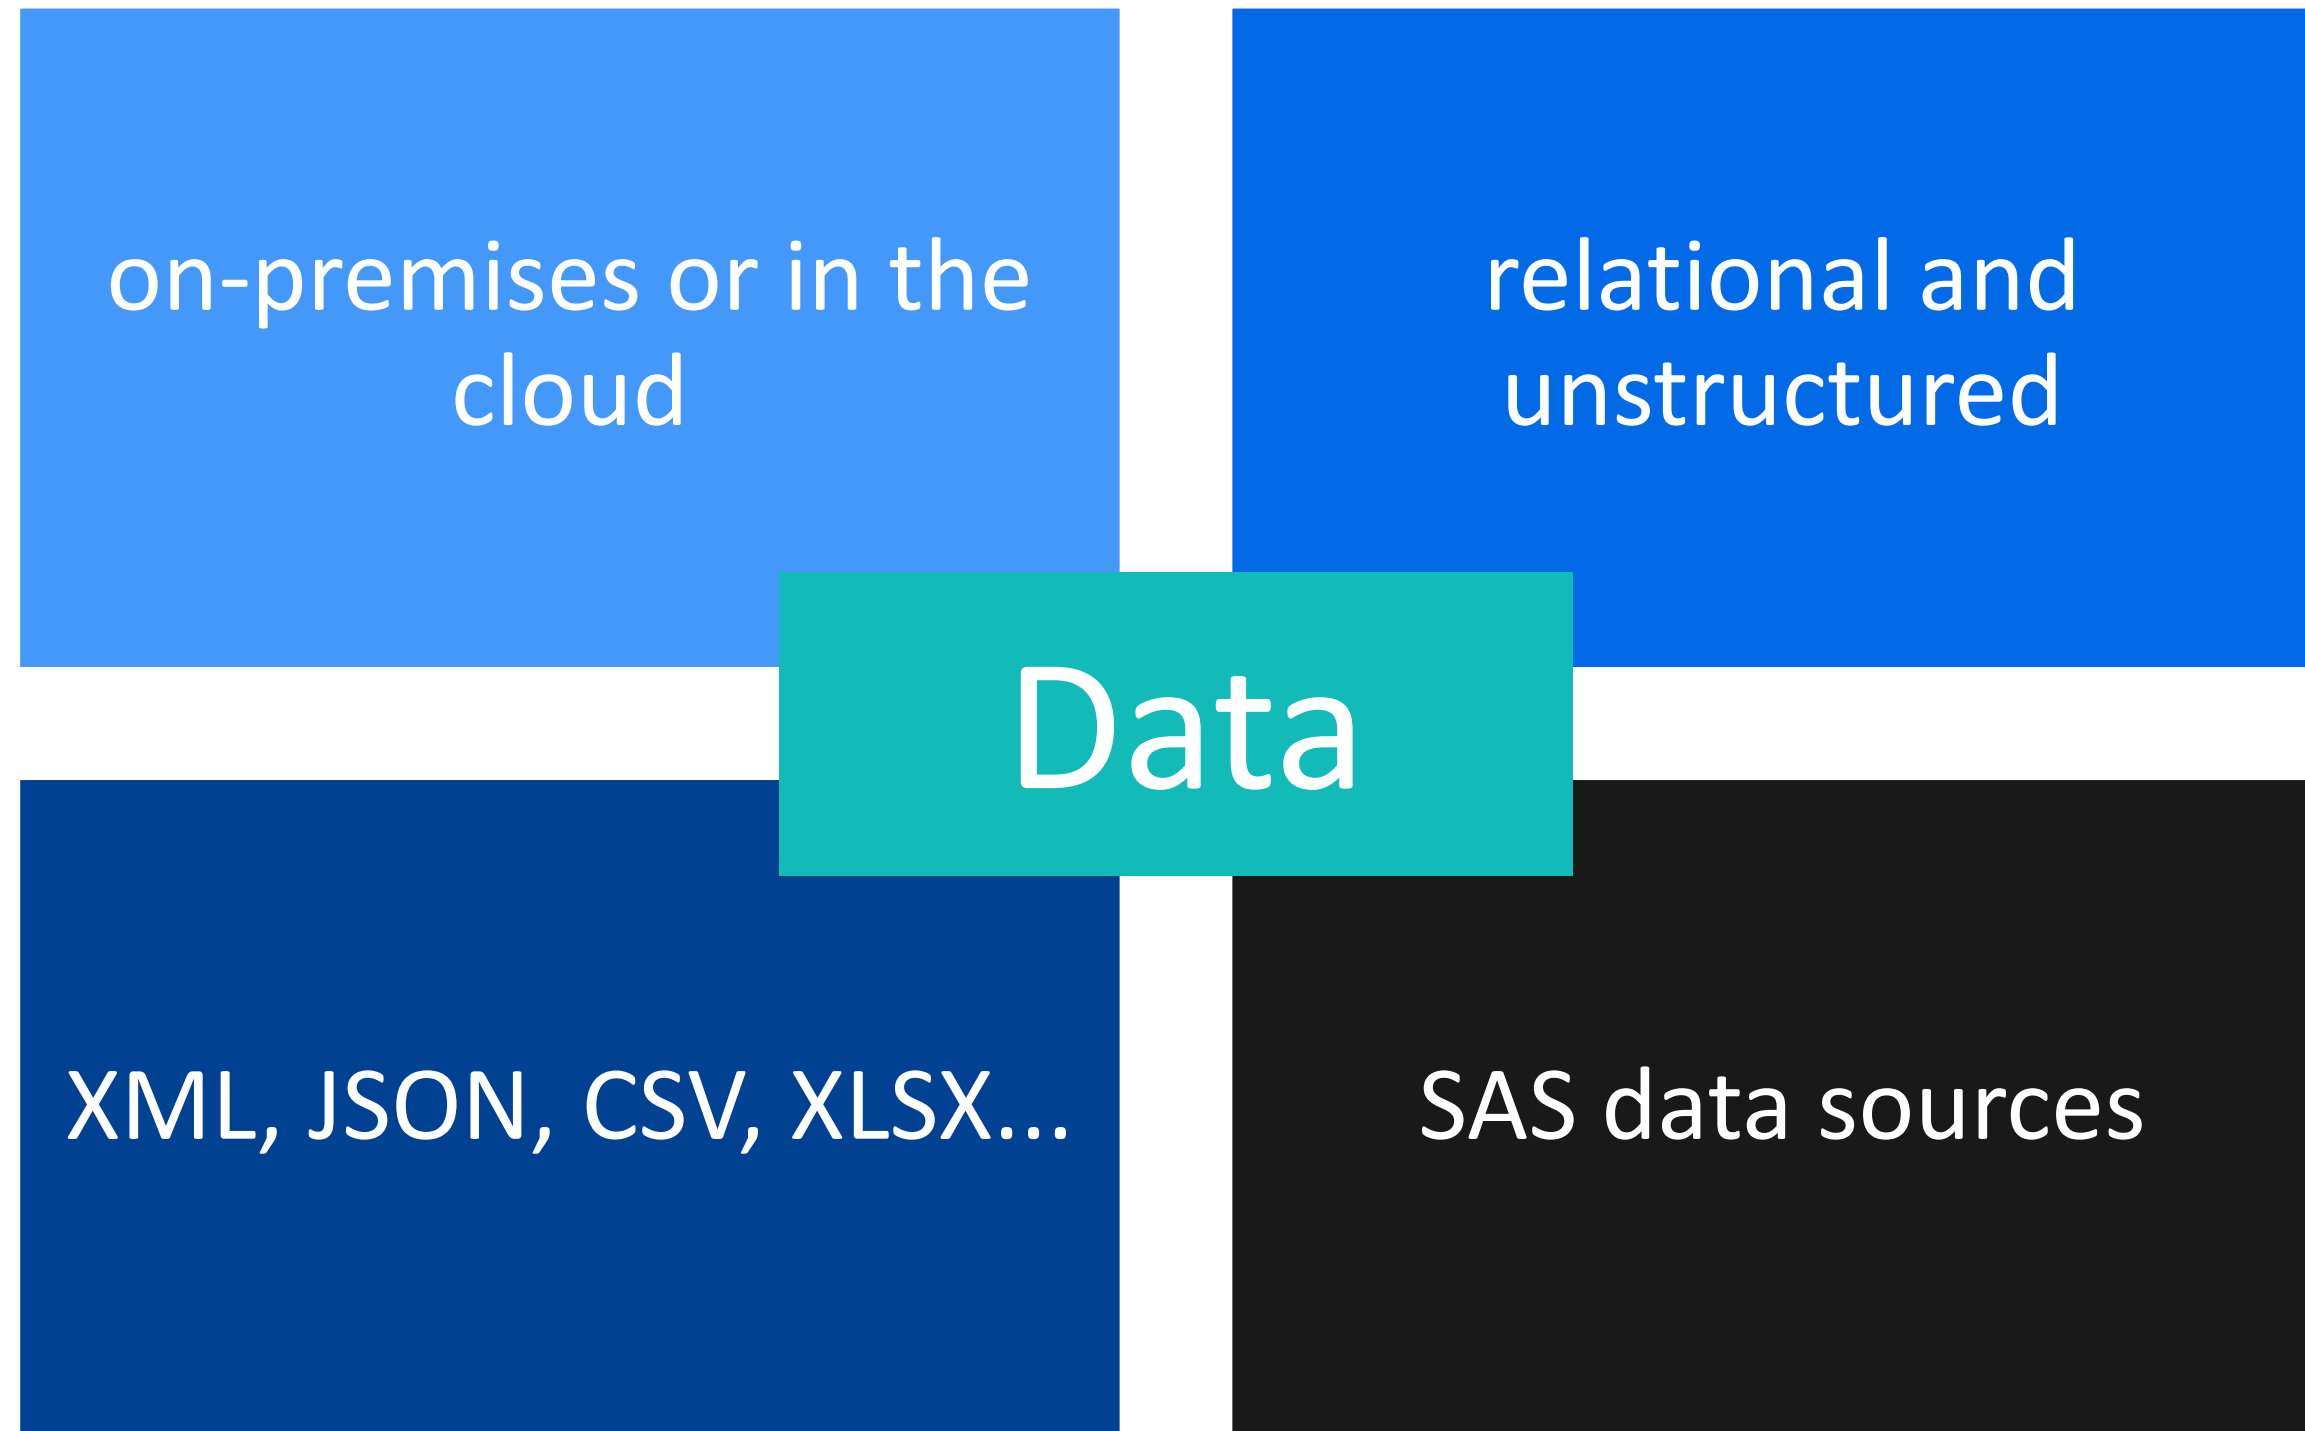
SAS® Applications

- SAS Studio (3.x)
- SAS Enterprise Guide
- SAS windowing environment

Other familiar applications can also be used to develop and submit code to SAS Viya.

Open Source


4 Data




5 Code

SAS Viya Programming Process

Use the SAS Code That You Are Familiar with



SAS Data Step



Does the same data step syntax work for in-memory data in CAS?


Yes, we just need to consider how the data is processed, in either a single thread or multiple threads.

Base SAS DATA Step Processing

Single-Threaded Processing



Single-threaded processing reads data sequentially, one row at a time.



SAS Viya DATA Step Processing

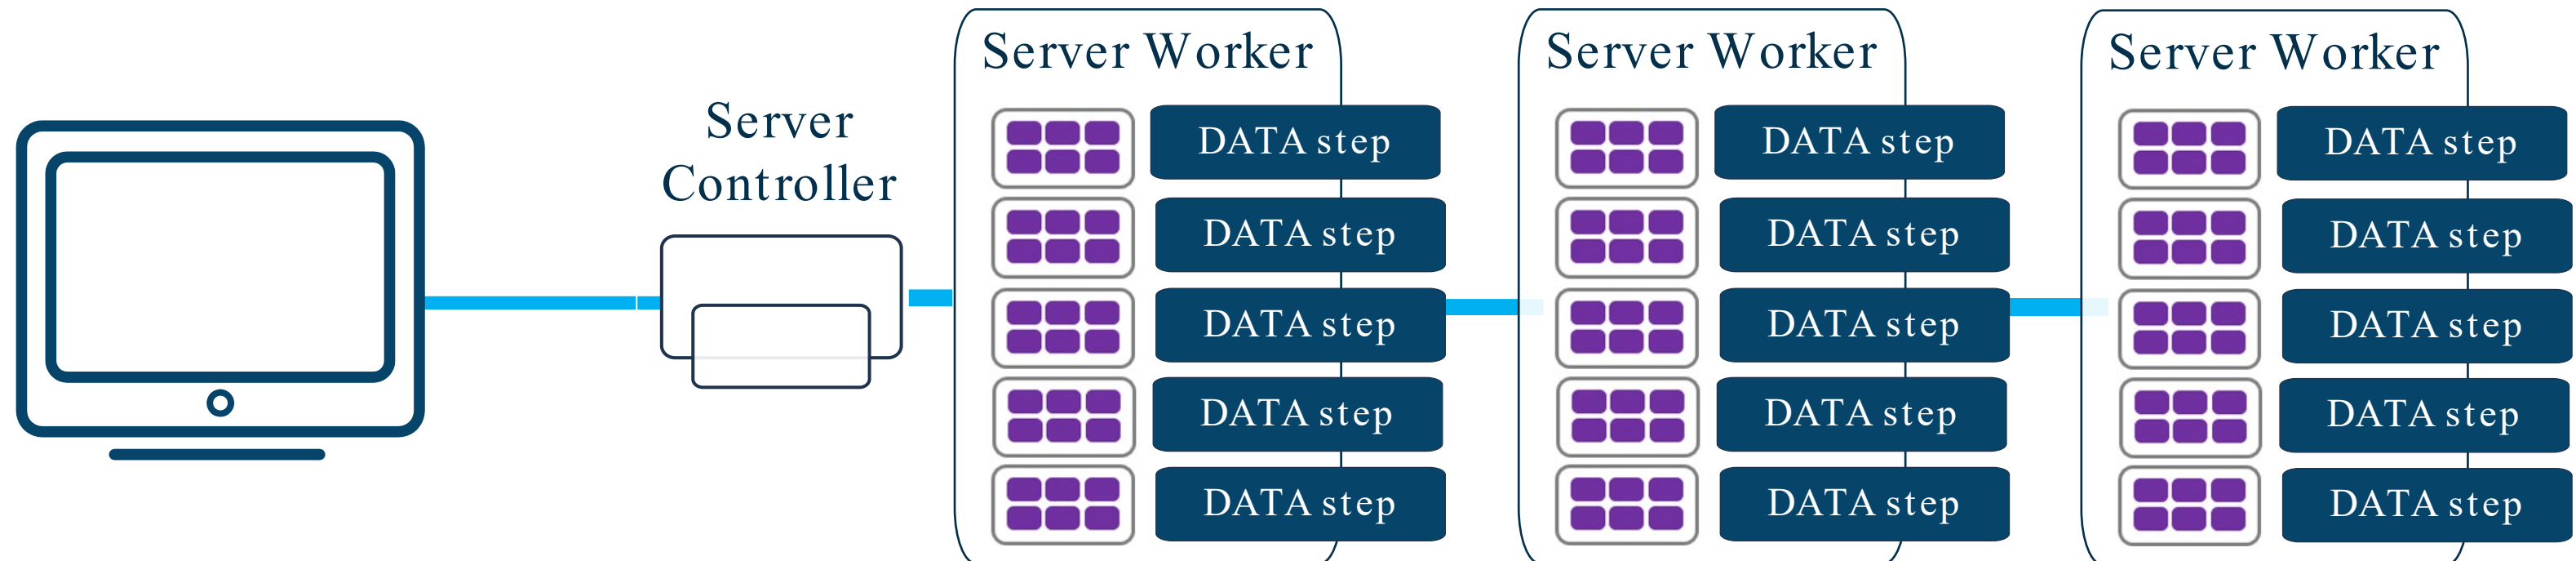
Massively Parallel Processing (MPP)



SAS Viya enables data to be divided and processed simultaneously on multiple threads.




SAS Cloud Analytic Services (CAS)




Let's Start With What We Know

```
LIBNAME libref engine "path";
```



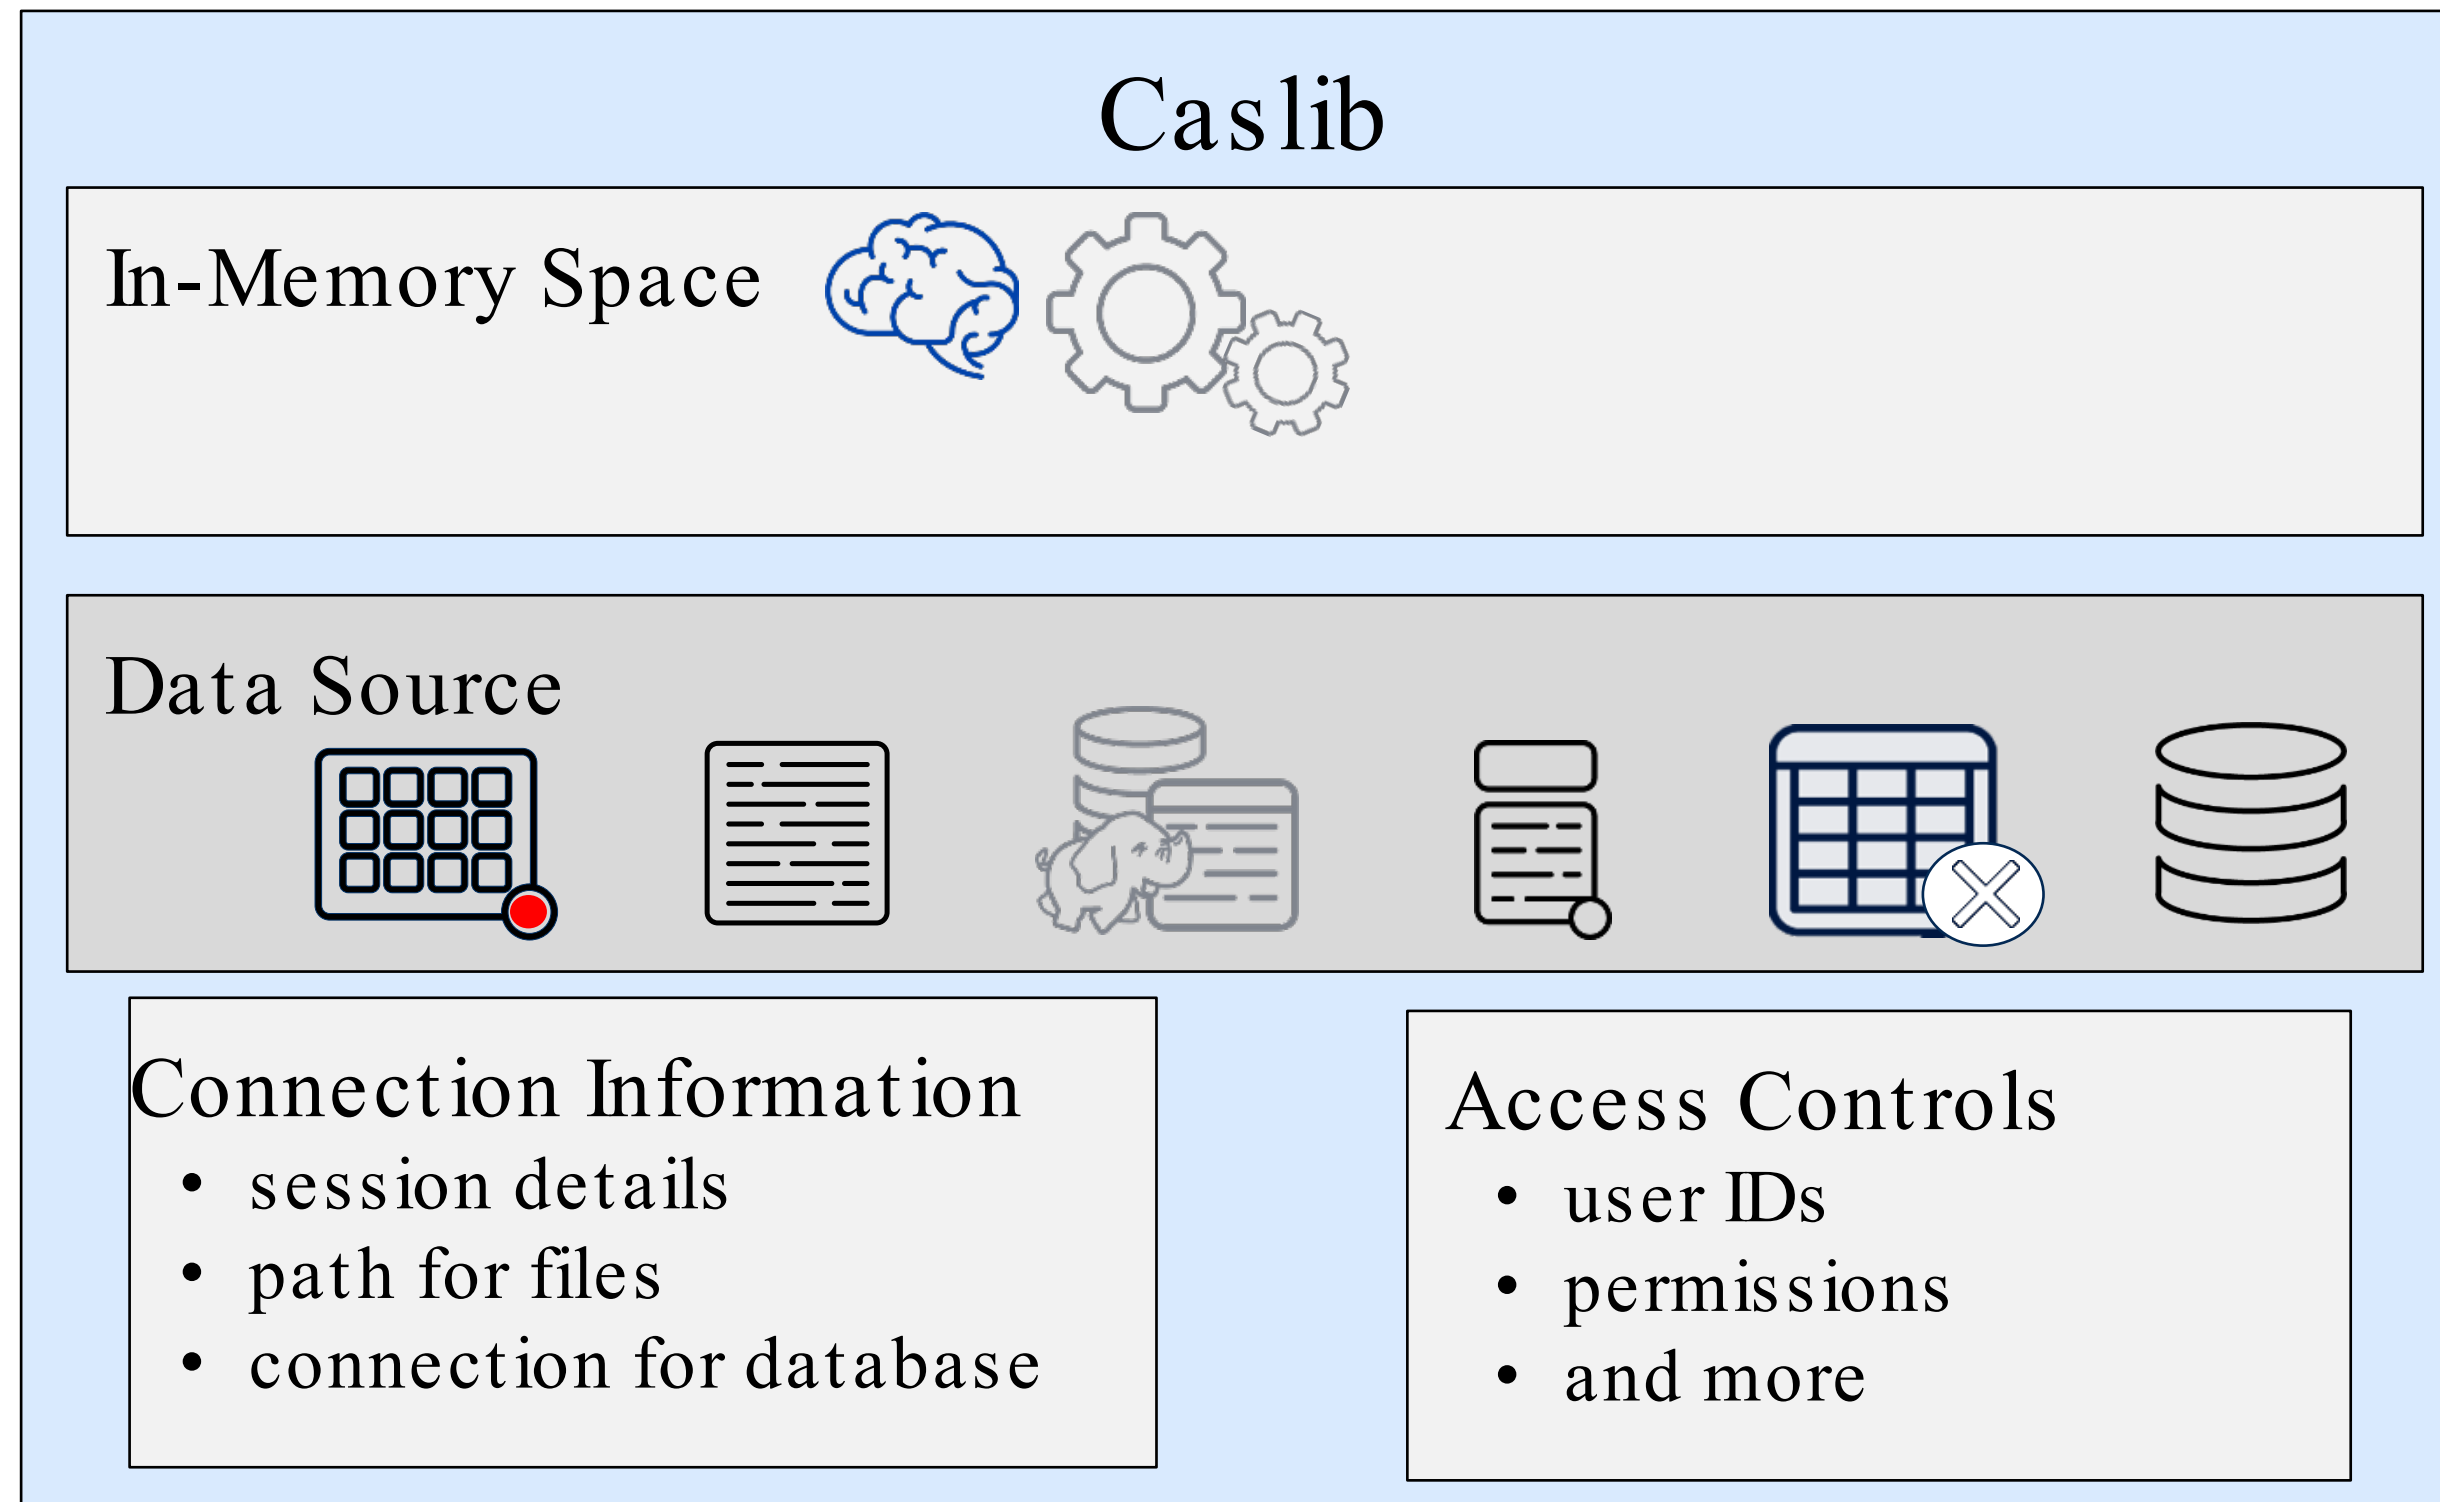
```
libname mysas "/workshop/pgvy34/data";
```



mysas

Caslibs


Used to Access Data in SAS Viya



Start A Session Via Code




```
cas mySession sessopts=( caslib=casuser timeout=1800 locale="en_US");
```



CAS session created is *mySession*

Sub option	Purpose
CASLIB=caslib	Specifies the name of the Caslib where the table is saved.
TIMEOUT=	Specifies the SAS Cloud Analytic Services session time-out in seconds for a new or existing session.
LOCALE=	Specifies a set of attributes in a SAS session that reflect the language, local conventions, and culture for a geographical region.

CAS-Enabled Steps



Data and Proc steps that are not CAS-enabled will automatically run on the Compute Server.

CAS-enabled steps are automatically converted to CAS actions behind the scenes. meaning SAS translated BASE SAS compliant code into CAS actions so your code can run in CAS/Viya)

```
proc casutil;  
    ...  
quit;  
  
data casuser.profit;  
    set casuser.orders;  
    ...  
run;  
  
proc means data=casuser.profit;  
    ...  
run;
```

DATA Step

```
DATA output-table;  
    . . .  
RUN;
```

modify existing values

compute new variables

conditionally process

combine tables

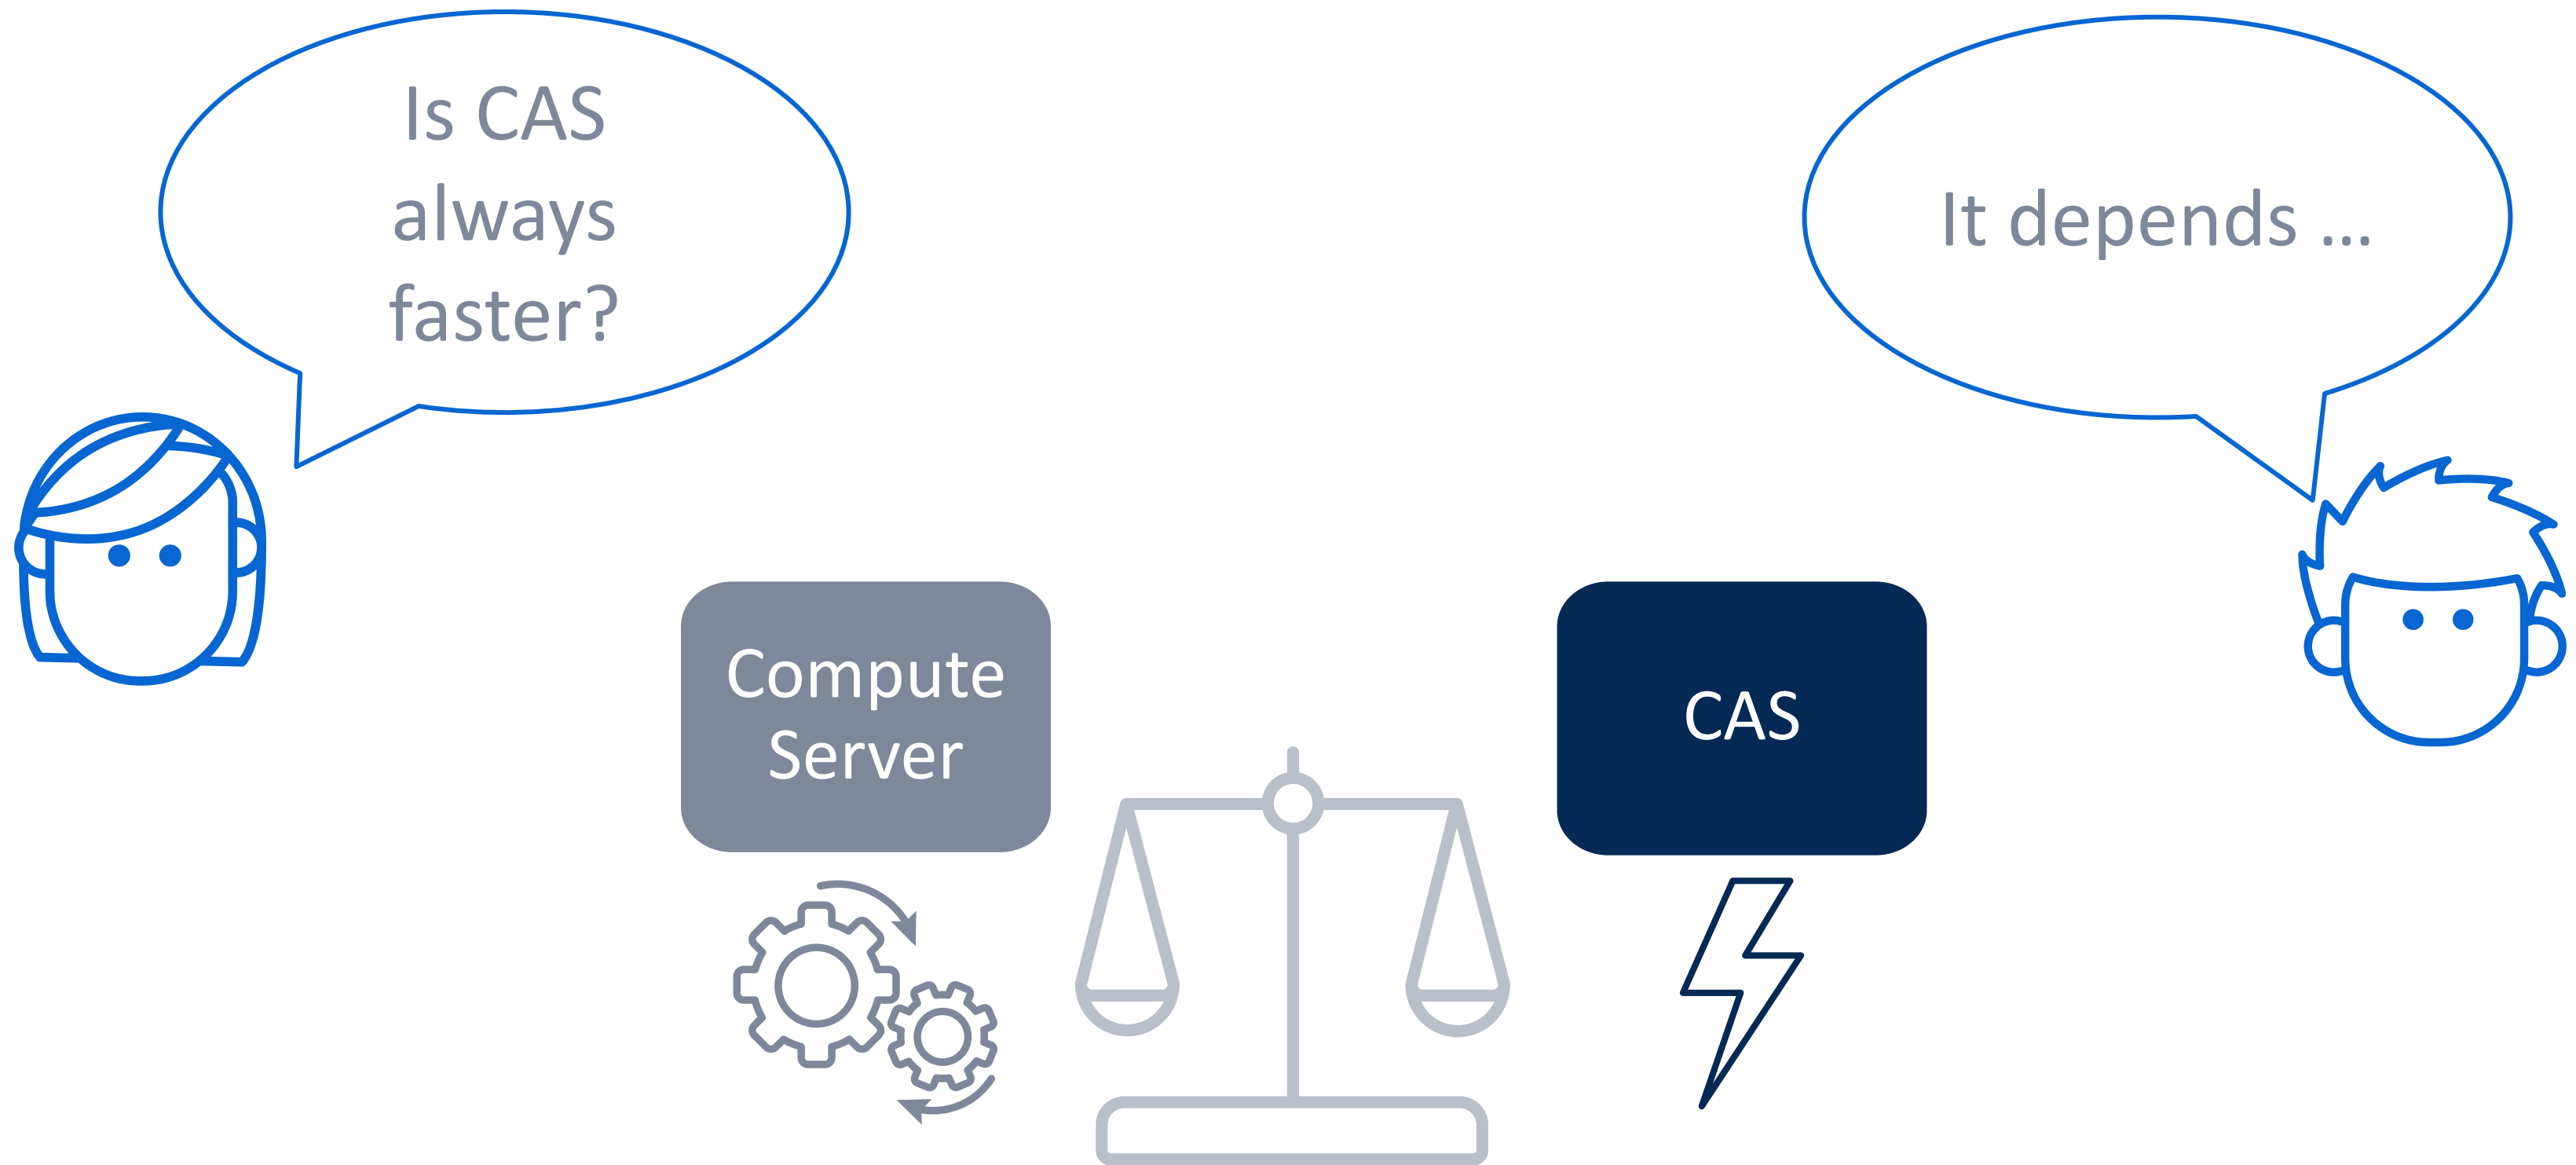
The DATA step is
used to manage
and manipulate
tables in
preparation for
analysis.


Modifying the DATA Step to Run in CAS

```
data casuser.Departments;  
  set casuser.mycustomers end=eof;  
  select(Continent);  
    when ('Africa', 'Asia')  
      Department="General Shoes";  
    when ("Oceania")  
      Department="Men's Shoes";  
    when ("North America", "Europe")  
      Department="Women's Shoes";  
    otherwise Department='Unknown';  
  end;  
  keep City Continent Department;  
  if eof then put _threadid_ _N_ =;  
run;
```


If the output table and input table use a caslib library reference, the DATA step runs in CAS.

Compute Server or CAS?







Start services



Load data into memory



Distribute data and processing
among worker nodes




Return results to controller



Deliver combined results to the client




Which steps
should I
modify to run
in CAS?



steps using data sources
larger than 50 GB



long-running steps (30 minutes+)



computationally demanding PROCs



DATA steps with many computations,
functions, or conditional logic

Handy Resources

[SAS Viya Programming](#)

[SAS Viya Learning Path](#)

[Data Step Processing in CAS](#)

[SAS Education Resources for Viya](#)

[SAS 9 and SAS Viya Functional Comparison](#)

[SAS Global Forum Tutorial - Coding in SAS Viya](#)

[Understanding SAS: The Different Processing Engines](#)

[Seriously Serial or Perfectly Parallel Data Transfer with SAS Viya](#)

Demo

Quiz Question 1

Programs written for SAS 9 must be modified before they can run in Viya.

TRUE/ FALSE

Quiz Answer 1

Programs written for SAS 9 must be modified before they can run in Viya.

FALSE

The SAS Viya Compute Server runs programs written for SAS 9.

Quiz Question 2

SAS Viya includes two compute engines called the Compute Server and _____.

Quiz Answer 2

SAS Viya includes two compute engines called the Compute Server and Cloud Analytic Services.

The Cloud Analytic Services massively parallel in-memory compute engine is often referred to as CAS.

Quiz Question 3

Which application enables you to prepare CAS tables and create CAS tables using transforms?

- a. SAS Studio
- b. SAS Data Explorer
- c. SAS Data Studio
- d. SAS Prepare Data

Quiz Answer 3

Which application enables you to prepare CAS tables and create CAS tables using transforms?

- a. SAS Studio
- b. SAS Data Explorer
- c. SAS Data Studio
- d. SAS Prepare Data

Quiz Question 4

What is the name of the high-performance, in-memory analytics engine used in SAS Viya?

- a. BAS (Business Analytic Services)
- b. CAS (Cloud Analytic Services)
- c. LAS (LASR Analytic Services)
- d. SAS (Statistic Analytic Services)

Quiz Answer 4

What is the name of the high-performance, in-memory analytics engine used in SAS Viya?

- a. BAS (Business Analytic Services)
- b. CAS (Cloud Analytic Services)
- c. LAS (LASR Analytic Services)
- d. SAS (Statistic Analytic Services)

Quiz Question 5

Which application enables users to create reports?

- a. SAS Report Creator
- b. SAS Visual Analytics
- c. SAS Report Viewer
- d. SAS Visual Analytics App

Quiz Answer 5

Which application enables users to create reports?

- a. SAS Report Creator
- b. SAS Visual Analytics
- c. SAS Report Viewer
- d. SAS Visual Analytics App

Quiz Question 6

What applications are new in the SAS Viya platform?

- a. SAS Studio
- b. SAS Metadata Server
- c. SAS Business Intelligence Platform
- d. SAS Access to PC Files

Quiz Answer 6

What applications are new in the SAS Viya platform?

- a. SAS Studio
- b. SAS Metadata Server
- c. SAS Business Intelligence Platform
- d. SAS Access to PC Files

For purposes of pure coding in SAS Viya, SAS Studio is a new application. SAS Metadata Server, SAS BI, SAS Access to PC Files are not new to SAS Viya

Quiz Question 7

What is the SAS compute server?

- a. The Base SAS 9 Engine with mostly single-threaded processing
- b. Cloud Analytics Server
- c. Multithreaded processing
- d. Massive parallel processing environment

Quiz Answer 7

What is the SAS compute server?

- a. The Base SAS 9 Engine with mostly single-threaded processing
- b. Cloud Analytics Server
- c. Multithreaded processing
- d. Massive parallel processing environment

The SAS Compute Server is the traditional SAS 9 engine that we are used to working with. Single-threaded processing for the most part, with a few thread enabled PROCs that support multithreaded processing Eg. MEANS, REPORT, SORT, SUMMARY, TABULATE, SQL

Quiz Question 8

You have been given the SAS Viya environment to work with. What happens to all the base SAS 9 data step code that you and your team spent months developing?

- a. With a few tweaks to your code, you can reuse your base SAS 9 code
- b. You have to rewrite all your code from scratch to make it SAS Viya compatible

Quiz Answer 8

You have been given the SAS Viya environment to work with. What happens to all the base SAS 9 data step code that you and your team spent months developing?

- a. With a few tweaks to your code, you can reuse your SAS 9 code
- b. You have to rewrite all your code from scratch to make it SAS Viya compatible

With just a few tweaks, to your Base SAS 9 code you can advantage of all the advantages of speed, in memory engines, multithreaded capacity and so much more in SAS Viya

Quiz Question 9

Which step should you modify to run in CAS?

- a. computationally demanding PROCs
- b. steps using data sources < 50 GB
- c. Short steps under 30 minutes
- d. simple data steps

Quiz Answer 9

Which step should you modify to run in CAS?

- a. computationally demanding PROCS
- b. steps using data sources < 50 GB
- c. Short steps under 30 minutes
- d. simple data steps

As you work in your Viya environment and consider writing or modifying programs to run in CAS, there are a few best practices to consider. CAS generally outperforms the Compute Server if you have data sources larger than 50 GB, steps that run for 30 minutes or longer, PROCs that are computationally demanding, or DATA steps with extremely long or complex logic.

Thank You

Charu Shankar
SAS Institute Toronto

EMAIL	Charu.shankar@sas.com
BLOG	https://blogs.sas.com/content/author/charushankar/
TWITTER	CharuYogaCan
LINKEDIN	https://www.linkedin.com/in/charushankar/

✓ Did you
enjoy this
session, Let us
know in the
[evaluation](#)

